

EMOTION RECOGNITION USING SPEECH PROCESSING

A PROJECT REPORT

Submitted by

JIRILIN BABU S [211421104114]

RENGANATHAN V [211421104206]

ROGAN J [211421104210]

*in partial fulfillment for the award of the degree
of*

BACHELOR OF ENGINEERING

IN

COMPUTER SCIENCE AND ENGINEERING



PANIMALAR ENGINEERING COLLEGE

(An Autonomous Institution, Affiliated to Anna University, Chennai)

APRIL 2025

PANIMALAR ENGINEERING COLLEGE

(An Autonomous Institution, Affiliated to Anna University, Chennai)

BONAFIDE CERTIFICATE

Certified that this project report “**EMOTION RECOGNITION USING SPEECH PROCESSING**” is the bonafide work of “**JIRILIN BABU S [211421104114], RENGANATHAN V [211421104206] and ROGAN J [211421104210]**” who carried out the project work under my supervision.

SIGNATURE

**Dr. L. JABASHEELA, M.E., Ph.D.,
HEAD OF THE DEPARTMENT,
PROFESSOR,**

Department of Computer Science and
Engineering,
Panimalar Engineering College,
Nazarathpet,
Poonamallee,
Chennai, 600 123.

SIGNATURE

**Dr. V. SATHIYA, M.E., Ph.D.,
SUPERVISOR,
PROFESSOR,**

Department of Computer Science and
Engineering,
Panimalar Engineering College,
Nazarathpet,
Poonamallee,
Chennai, 600 123.

Certified that the above candidate(s) was examined in the End Semester Project
Viva-Voce Examination held on

INTERNAL EXAMINER

EXTERNAL EXAMINER

DECLARATION BY THE STUDENT

We, **JIRILIN BABU S [211421104114]**, **RENGANATHAN V [211421104206]** and **ROGAN J [211421104210]** hereby declare that this project report titled “**EMOTION RECOGNITION USING SPEECH PROCESSING**”, under the guidance of **Dr. V. SATHIYA** is the original work done by us and we have not plagiarized or submitted to any other degree in any university by us.

JIRILIN BABU S [211421104114]

RENGANATHAN V [211421104206]

ROGAN J [211421104210]

ACKNOWLEDGEMENT

Our profound gratitude is directed towards our esteemed Secretary and Correspondent, **Dr. P. CHINNADURAI, M.A., Ph.D.**, for his fervent encouragement. His inspirational support proved instrumental in galvanizing our efforts, ultimately contributing significantly to the successful completion of this project.

We want to express our deep gratitude to our Directors, **Tmt. C. VIJAYARAJESWARI, Dr. C. SAKTHI KUMAR, M.E., Ph.D.**, and **Dr. SARANYASREE SAKTHI KUMAR, B.E., M.B.A., Ph.D.**, for graciously affording us the essential resources and facilities for undertaking of this project.

Our gratitude is also extended to our Principal, **Dr. K. MANI, M.E., Ph.D.**, whose facilitation proved pivotal in the successful completion of this project.

We express our heartfelt thanks to **Dr. L. JABASHEELA, M.E., Ph.D.**, Head of the Department of Computer Science and Engineering, for granting the necessary facilities that contributed to the timely and successful completion of project.

We would like to express our sincere thanks to Project Coordinator **Dr. R. JOSPHINELEELA, M.E., Ph.D.**, and Project Guide **Dr. SATHIYA, M.E., Ph.D.**, and all the faculty members of the Department of CSE for their unwavering support for the successful completion of the project.

JIRILIN BABU S [211421104114]
RENGANATHAN V [211421104206]
ROGAN J [211421104210]

ABSTRACT

Emotion is an essential part of human communication, playing a critical role in conveying intentions, mood, and underlying meaning beyond words. Understanding these emotions from speech is valuable across diverse fields from enhancing healthcare systems (e.g., monitoring patients' emotional states) to improving virtual assistants (e.g., Alexa, Siri) and revolutionizing human-computer interaction (e.g., making chatbots more empathetic and responsive). However, achieving accurate Speech Emotion Recognition (SER) is challenging due to variations in speech tone, language, background noise, and speaker diversity. This project proposes a Hybrid Neural Fusion (HNF) based Speech Emotion Recognition (SER) system that addresses these challenges by combining Convolutional Neural Networks (CNN) and Long Short-Term Memory (LSTM) networks. This hybrid architecture leverages the strengths of both models – CNNs are excellent at extracting spatial features from audio representations (like spectrograms), while LSTMs are designed to capture temporal patterns, making them ideal for analyzing how emotions evolve over time.

LIST OF TABLES

TABLE NO.	TABLE NAME	PAGE NO.
1	Confusion Matrix	21

LIST OF FIGURES

FIGURE NO.	FIGURE NAME	PAGE NO.
1	System Architecture	11
2	Training the Model	49
3	Server	50
4	Selecting an Audio	50
5	Tracked Emotion	51
6	Confusion Matrix	51

LIST OF SYMBOLS

SYMBOL NO.	SYMBOL	SYMBOL NAME
1	SER	Speech Emotion Recognition
2	CNN	Convolutional Neural Network
3	LSTM	Long Short-Term Memory
4	AI	Artificial Intelligence
5	MFCC	Mel Frequency Cepstral Coefficient
6	ERC	Emotion Recognition in Conversation
7	DAN	Distribution Alignment Network
8	IDA	Implicit Distribution Alignment
9	HMM	Hidden Markov Model

TABLE OF CONTENTS

CHAPTER NO.	TITLE	PAGE NO.
	ABSTRACT	v
	LIST OF TABLES	vi
	LIST OF FIGURES	vii
	LIST OF SYMBOLS	viii
1.	INTRODUCTION	1
	1.1. Problem Definition	2
2.	LITERATURE REVIEW	4
3.	THEORETICAL BACKGROUND	9
	3.1. Implementation Environment	9
	3.2. System Architecture	11
	3.2.1. Input Handling	11
	3.2.2. Preprocessing	12
	3.2.3. Feature Extraction	12
	3.2.4. Emotion Classification	13
	3.2.5. Performance Visualization	14
4.	SYSTEM IMPLEMENTATION	16
	4.1. Proposed Model	16
	4.1.1. Preprocessing	16
	4.1.2. Feature Extraction	16
	4.1.3. Model Architecture	17
	4.1.4. Performance Evaluation	18
5.	RESULTS & DISCUSSION	24
	5.1. Results	24

6.	CONCLUSION	27
6.1.	Future Enhancements	28
	APPENDICES	30
A.1	SDG Goals	30
A.2	Source Code	33
A.3	Screenshots	49
A.4	Plagiarism Report	52
	REFERENCES	65

CHAPTER 1

INTRODUCTION

1. INTRODUCTION

Emotion detection from speech is emerging as a pivotal area in Artificial Intelligence (AI) and machine learning, driven by the need for machines to understand not just what we say but how we say it. Human communication is a rich, multi-layered process that conveys meaning not only through words but also through non-verbal cues such as tone, pitch, rhythm, and intensity. These vocal characteristics reflect underlying emotions, helping listeners interpret whether a statement is happy, sad, angry, or neutral. Traditional AI systems, however, focus primarily on text-based analysis, overlooking these vital vocal cues leading to responses that feel mechanical and disconnected from human emotions.

Speech Emotion Recognition (SER) systems aim to bridge this gap by equipping machines with the ability to perceive and interpret emotions from speech transforming them from passive responders to empathetic communicators. By analyzing the acoustic properties of speech, such as pitch variations (e.g., high pitch for excitement or fear), rhythm patterns (e.g., slow, drawn-out speech for sadness), and vocal intensity (e.g., sharp, loud tones for anger), SER systems can decode the speaker's emotional state enhancing machine interactions with humans.

The potential applications are vast and transformative. Virtual assistants like Alexa, Siri, and Google Assistant can become more emotionally aware, responding compassionately to frustration or enthusiastically to excitement. Call centers equipped with SER technology can detect customer dissatisfaction in real-time, prompting agents to adjust their approach reducing escalations and improving customer satisfaction. In mental health applications, SER systems can track emotional patterns from daily speech, providing early indicators of emotional distress empowering users and healthcare professionals to intervene sooner.

As the technology evolves, emotion detection from speech holds the potential to redefine human-computer interaction, making machines more intuitive, empathetic, and human-like fostering a future where AI understands not just what we say, but how we feel when we say it.

1.1. Problem Definition

Existing Speech Emotion Recognition (SER) systems face significant limitations that hinder their real-world applicability. These systems, which perform well under controlled, noise-free conditions, often fail to deliver accurate results when deployed in diverse environments due to challenges like background noise, speaker variability, and shallow feature extraction.

Real-world audio environments are rarely clean, frequently featuring background noise such as environmental sounds, overlapping conversations, or microphone distortions. These interferences obscure the emotional content of speech, and traditional SER models trained on clean datasets struggle to distinguish relevant emotional cues from irrelevant noise. For example, sadness conveyed in a noisy street might be misclassified as neutrality or anger due to interference from background sounds.

Additionally, speaker variability presents a critical obstacle. Human speech differs widely across individuals due to variations in pitch, tone, accents, and speaking styles. Current SER systems, often trained on datasets with limited speaker diversity, exhibit speaker dependency, performing well on familiar voices but failing to generalize to new, unseen speakers. This restricts the system's adaptability, particularly in applications like virtual assistants or customer service, which demand robust performance across diverse voices.

CHAPTER 2

LITERATURE REVIEW

2. LITERATURE REVIEW

Project Title : Speech Emotion Recognition Using Synthetic Data Augmentation from Emotion Conversion

Author Name : Karim M. Ibrahim, Antony Perzo, Simon Leglaive

Year of Publish : 2024

Abstract: One of the main challenges in speech emotion recognition is the lack of large 4abeled datasets. The progress in speech synthesis allows us to generate reliable and realistic expressive speech. In this work, we propose using a state-of-the-art end-to-end speech emotion conversion model to generate new synthetic data for training speech emotion recognition models. We first evaluate the quality of the converted speech on new unseen datasets, which proves to be on par with the training data. Then, we study the effect of using the synthesized speech as data augmentation. We show that this approach improves the overall performance of emotion recognition models on two different datasets, IEMOCAP and RAVDESS, both in the cases of speaker dependent and independent emotion recognition using a fine-tuned wav2vec 2.0.

Project Title : Speech Emotion Recognition in Conversation Using Graph Convolutional Networks

Author Name : Deeksha Chandola, Enas Altarawneh

Year of Publish : 2024

Abstract: Speech emotion recognition (SER) is the task of automatically recognizing emotions expressed in spoken language. Current approaches focus on analyzing isolated speech segments to identify a speaker’s emotional state. Meanwhile, recent text-based emotion recognition methods have effectively shifted towards emotion recognition in conversation (ERC) that considers conversational context. Motivated by this shift, here we propose SERC-GCN, a method for speech emotion recognition in conversation (SERC) that predicts a

speaker’s emotional state by incorporating conversational context, speaker interactions, and temporal dependencies between utterances. SERC-GCN is a two-stage method. First, emotional features of utterance-level speech signals are extracted. Then, these features are used to form conversation graphs that are used to train a graph convolutional network to perform SERC. We empirically evaluate the effectiveness of SERC-GCN and show that it outperforms the current state-of-the-art methods on the IEMOCAP benchmark dataset.

Project Title : Deep Implicit Distribution Alignment Networks for cross-Corpus Speech Emotion Recognition

Author Name : Yan Zhao

Year of Publish : 2023

Abstract: A novel deep transfer learning method called deep implicit distribution alignment networks (DIDAN) to deal with cross-corpus speech emotion recognition (SER) problem, in which the labeled training (source) and unlabeled testing (target) speech signals come from different corpora. Specifically, DIDAN first adopts a simple deep regression network consisting of a set of convolutional and fully connected layers to directly regress the source speech spectrums into the emotional labels such that the proposed DIDAN can own the emotion discriminative ability. Then, such ability is transferred to be also applicable to the target speech samples regardless of corpus variance by resorting to a well-designed regularization term called implicit distribution alignment (IDA). Unlike widely-used maximum mean discrepancy (MMD) and its variants, the proposed IDA absorbs the idea of sample reconstruction to implicitly align the distribution gap, which enables DIDAN to learn both emotion discriminative and corpus invariant features from speech spectrums. To evaluate the proposed DIDAN, extensive cross-corpus SER experiments on widely-used speech emotion corpora are carried out.

Experimental results show that the proposed DIDAN can outperform lots of recent state-of-the-art methods in coping with the cross-corpus SER tasks.

Project Title : Machine Learning for Speech Emotion Recognition

Author Name : Darwindra, Sofia Sa'Idah

Year of Publish : 2023

Abstract: Speech Emotion Recognition (SER) has been under the scrutiny of researchers for many years. The emotion recognition system through speech is the ability to identify the type of emotion from a human speech to create a natural interaction between humans and machines (computers). This research uses the Linear Predictive Coding (LPC) feature extraction method with orders of 1, 8, and 16. Then the classification method used is the Hidden Markov Model (HMM) and K-Nearest Neighbor (KNN) methods. This research aims to test the LPC method in extracting features without going through the frame blocking and windowing processes and to compare the classification results between KNN and HMM. The use of these two methods is also to test how well the LPC feature extraction results are for the development of SER. Based on the testing process carried out, the best system performance in identifying the type of emotion is when using the LPC parameter of order 16. The average level of accuracy obtained is 96.88% for the KNN classification with the number of neighbors 1, 3, and 5. Meanwhile, when using HMM classification, the highest level of accuracy obtained is 62.50% with the number of states 10.

Project Title : Exploring Fusion Techniques for Multimodal Emotion Recognition

Author Name : Romania, Horia Cucu

Year of Publish : 2024

Abstract: Emotion recognition is crucial for improving human-computer interaction, allowing machines to understand and respond to emotions. This

enhances user experience and personalizes interactions, benefiting applications like virtual assistants and mental health diagnostics. In this paper, we propose a multimodal emotion recognition system that relies on speech and visual information. For the speech-based modality, we conducted experiments using a CNN architecture along with five types of handcrafted speech features. For visual emotion recognition, we evaluated a ResNet18-based network, leveraging its potential by fine-tuning a pretrained model based on a larger dataset, ImageNet. After developing these two separate classifier architectures, one for speech and the other for visual information, we studied multiple late fusion techniques to effectively maximize emotion recognition from these two sources of information. Our best multimodal system achieved an accuracy of 97.95%, demonstrating an improvement over the individual modalities and other benchmark systems.

CHAPTER 3

THEORETICAL BACKGROUND

3. THEORETICAL BACKGROUND

3.1. IMPLEMENTATION ENVIRONMENT

The implementation environment for this Speech Emotion Recognition (SER) project consists of a carefully chosen set of tools, libraries, and frameworks to ensure smooth development, training, and deployment. Each component plays a crucial role in creating an efficient, robust, and interactive SER system that processes audio, extracts meaningful features, trains a hybrid model, and visualizes performance metrics in real time.

3.1.1. Programming Language and Frameworks

The project is built using Python, thanks to its versatility, rich ecosystem of libraries, and strong support for machine learning and audio processing. The core machine learning framework is PyTorch – a powerful deep learning library that enables flexible model architecture design, efficient tensor computations, and GPU acceleration for faster training. PyTorch’s dynamic computation graph makes it easier to experiment with different model architectures, particularly the Hybrid Neural Fusion (HNF) setup used in this project. For the web application component, Flask is employed. Flask is a lightweight, yet powerful, Python web framework that facilitates building interactive web dashboards and backend systems. It allows the SER model to be integrated with a user-friendly interface, enabling users to upload audio files, record live speech, and monitor performance metrics in real time.

3.1.2. Audio Processing and Feature Extraction Tools

Librosa, a Python library designed specifically for audio analysis, is used for loading audio files, extracting features like MFCCs, Chroma, Spectral Contrast, and Tonnetz, and handling sampling rates. This ensures that input audio is preprocessed accurately and consistently before being fed into the

model. Additionally, Pyaudio supports live recording directly from the microphone, making it possible to capture audio on-the-fly for real-time emotion predictions. Soundfile is used to save the recordings into a ``.wav`` format that the model can process immediately.

3.1.3. Model Development and Training

The hybrid neural fusion model combines a Convolutional Neural Network and LSTM network, implemented using PyTorch. The CNN extracts spatial features from spectrogram-like audio representations, while the LSTM learns temporal dependencies – essential for recognizing emotions that unfold over time. The training process involves optimizing the model using Adam optimizer with cross-entropy loss, ensuring efficient weight updates and faster convergence.

3.1.4. Performance Evaluation and Visualization

For performance evaluation, Scikit-learn provides a comprehensive suite of metrics, including Accuracy, Precision, Recall, F1 Score, and Confusion Matrix. The confusion matrix is visualized using Seaborn, generating heatmaps that reveal how well the model distinguishes between emotions. The web dashboard is enhanced with HTML, CSS, and JavaScript to create an engaging, real-time performance monitor. JavaScript fetches updated performance data every 2 seconds, ensuring users can continuously track system metrics without manual refreshes.

3.1.5. Hardware and Runtime Environment

The project runs on Windows/Linux environments with a Python 3.9 installation. For model training and heavy tensor computations, a GPU-enabled machine with CUDA support is recommended – particularly for faster training

cycles on large datasets. The deployment, however, is lightweight and can run on CPU efficiently for real-time predictions.

3.2. SYSTEM ARCHITECTURE

The system architecture for this Speech Emotion Recognition (SER) project is designed to efficiently process audio data, extract meaningful features, and classify emotions using a hybrid model, and present results through an interactive web interface. It consists of five major components: Input Handling, Preprocessing, Feature Extraction, Emotion Classification, and Performance Visualization – all seamlessly integrated to ensure real-time predictions and performance monitoring. The architecture diagram is illustrated below.




Figure 1: System Architecture

3.2.1. Input Handling

The system supports two primary methods for collecting audio data. Users can upload `.wav` audio files directly through the web interface. This allows pre-recorded samples to be analyzed. Users can record speech in real time using their microphone. The recording is captured as a `.wav` file, ensuring consistency with the file upload format. This flexibility ensures the system can work with both pre-existing audio datasets and live, spontaneous speech a crucial feature for practical deployment.

3.2.2. Preprocessing

Once the audio is received, it undergoes preprocessing to ensure uniformity across different recordings. This includes:

- Loading the audio at its original sample rate using the Librosa library.
- Resampling the audio to a consistent rate (e.g., 22050 Hz) for consistent feature extraction.
- Noise Reduction and Silence Trimming (optional enhancement): This ensures the model focuses on the actual speech content rather than background noise or silence.

The preprocessing stage prepares the audio data to extract accurate features, ensuring the hybrid model receives high-quality inputs.

3.2.3. Feature Extraction

The system extracts a rich set of audio descriptors that capture both the timbral and tonal characteristics of speech. These features are crucial for distinguishing emotions like happiness, sadness, or anger. The Speech Features include:

MFCCs (Mel Frequency Cepstral Coefficients): Captures the short-term power spectrum, simulating how humans perceive sound. It's the most important feature for speech-based tasks.

Chroma: Represents the harmonic content based on 12-tone pitch classes, useful for identifying tonal variations associated with different emotions.

Spectral Contrast: Measures the difference between peaks and valleys in the frequency spectrum, helping differentiate speech textures (e.g., tense vs. relaxed tones).

Tonnetz: Captures tonal and harmonic characteristics, helping distinguish subtle emotional variations (e.g., fear vs. surprise).

The extracted feature vectors are then normalized to ensure consistent data input across different recordings.

3.2.4. Emotion Classification (Hybrid Neural Fusion (HNF) Model)

The system's core is built around a Hybrid Neural Fusion (HNF) Model, which is specifically engineered to extract and combine spatial and temporal information from audio inputs. This combination allows for a nuanced analysis of emotional expressions present in audio data. By leveraging both CNN and LSTM architectures, the model captures intricate details that contribute to its highly effective emotion recognition capabilities. The CNN (Convolutional Neural Network) processes spatial patterns from audio features to identify local details such as pitch variations or bursts of energy. These are often associated with specific emotions, such as the sharpness observed in anger or the softness found in sadness. The CNN layers are adept at learning frequency-based characteristics that distinguish various emotional states, laying the groundwork for precise spatial analysis. On the other hand, the LSTM (Long Short-Term Memory) network focuses on time-dependent sequences, crucial for understanding how emotional states evolve over time. Emotions like fear or surprise frequently emerge through shifts in tone and pace, and the LSTM excels at capturing these temporal changes. By processing sequential data, it provides insights into the dynamic progression of emotions.

Finally, a fully connected layer integrates the spatial and temporal patterns extracted by the CNN and LSTM layers. This layer predicts the emotion, classifying it into one of the seven categories: Neutral, Happy, Sad, Angry, Fearful, Disgust, or Surprised. Alongside the predicted emotion, the model also outputs probabilities for each class, ensuring a comprehensive analysis of the audio input.

3.2.5. Performance Visualization

The performance visualization is aimed at ensuring transparency and user-friendliness for the emotions predicted. This stage allows users to view predictions and performance metrics instantly, facilitating a seamless experience. The Emotion Prediction Display offers an immediate visualization of the predicted emotion as soon as the audio file is uploaded or live recording is completed. This feature delivers instant feedback, enabling users to see the results without any delay. An Emotion Tracking Graph complements this by presenting a dynamic line graph that showcases fluctuations in emotional states across multiple predictions. This visualization provides valuable insights into how emotions change over time, fostering deeper understanding of trends and patterns in the data. In addition to these real-time visuals, the system computes key performance metrics such as Accuracy, Precision, Recall, and the F1 Score. These metrics are crucial for evaluating the model's overall effectiveness and reliability. They offer users a clear perspective on how well the model is performing in identifying and classifying emotions. The Confusion Matrix serves as another powerful visualization tool, displayed in the form of a heatmap. It reveals detailed information about the model's prediction performance for each emotion class, shedding light on common misclassifications. For instance, the matrix can highlight patterns like the confusion between similar emotions, such as fear and surprise. To ensure that these metrics remain up-to-date, the dashboard is designed to refresh automatically every 2 seconds. This feature keeps the visualizations current as new predictions are made, reinforcing the system's efficiency and responsiveness.

CHAPTER 4

SYSTEM IMPLEMENTATION

4. SYSTEM IMPLEMENTATION

4.1. PROPOSED MODEL

The project is divided into four modules, each constituting specific features and procedures in detecting the emotion given. The modules are as follows:

4.1.1. Preprocessing

The preprocessing module is dedicated to acquiring audio data from a wide range of sources and meticulously preparing it for the feature extraction process. It begins by collecting speech samples that cover an extensive array of emotions, including neutral, happy, sad, angry, fearful, disgust, and surprised, ensuring a comprehensive dataset for emotion recognition. The preprocessing phase standardizes the audio data by resampling it to a consistent sampling rate, such as 16 kHz, which promotes uniformity across all inputs. Additionally, noise reduction techniques, like spectral gating and band-pass filtering, are applied to eliminate background noise and enhance the clarity of the recordings. To further enrich the dataset and improve the model's ability to generalize to real-world scenarios, data augmentation methods such as time stretching, pitch shifting, and the addition of background noise are employed. These augmentation strategies not only diversify the dataset but also bolster the model's robustness, making it more adept at identifying emotions in varied and unpredictable environments.

4.1.2. Feature Extraction

The feature extraction module is dedicated to transforming raw audio data into a structured numerical representation that can be effectively utilized for model training. At its core, it leverages various audio features to encapsulate

key aspects of speech, enhancing the model's ability to recognize emotions. MFCCs (Mel Frequency Cepstral Coefficients) are employed to capture the timbral, or tone color, properties of speech, mimicking the way human ears perceive sound by analyzing frequency content. Chroma features focus on the pitch content, mapping audio signals to the 12-tone musical scale, which is particularly valuable for detecting patterns in speech melody. Spectral Contrast highlights the differences between peaks and valleys in the audio spectrum, distinguishing between high-energy elements (such as those in anger) and low-energy elements (like sadness). Furthermore, Tonnetz identifies harmonic and tonal relationships within the signal, making it instrumental in differentiating emotional tones, such as the somber characteristics of sadness versus the intensity of anger. Collectively, these features provide a comprehensive, multi-dimensional input to the model, ensuring that emotions are captured through both voice tone and melodic variations in speech, ultimately enriching the emotional recognition process.

4.1.3. Model Architecture

The model architecture delves deeply into the Hybrid Neural Fusion (HNF) architecture, which combines the strengths of both Convolutional Neural Networks (CNNs) and Long Short-Term Memory (LSTM) networks to effectively recognize emotions in audio data. The CNN layers focus on extracting spatial patterns from audio spectrograms, learning intricate local features such as frequency peaks and energy bursts that are critical for distinguishing between tonal emotions. For example, these features help identify the sharpness of anger or the softness of sadness. Complementing this, the LSTM layers handle temporal dependencies by tracking how emotions evolve over time, which is especially important for understanding dynamic transitions in audio, such as a shift from a neutral tone to anger. Once the spatial and temporal patterns are extracted, a Fully Connected Layer takes on the role

of the decision-making mechanism, utilizing a Softmax activation function to classify the input into distinct emotional categories. To ensure effective training, the model employs cross-entropy loss, an appropriate choice for multi-class classification tasks, and optimizes performance using the Adam optimizer with a low learning rate of 0.0001 to achieve stable and consistent convergence. Furthermore, data augmentation techniques such as time stretching, pitch shifting, and background noise addition enhance the dataset's diversity. This step is vital for making the model robust against variations in speaker attributes and noisy environments, thereby increasing its ability to generalize effectively to real-world scenarios.

4.1.4. Performance Evaluation

The performance evaluation module is dedicated to evaluating the model's performance, using key metrics to assess its accuracy and reliability. The performance parameters to evaluate the model's performance are given below.

Accuracy

Accuracy is a metric used to measure the overall correctness of a model's predictions. It calculates the ratio of correctly predicted instances to the total number of predictions made. Simply put, accuracy helps us understand how well the model performs in terms of correctly identifying emotions across the dataset. The formula for calculating accuracy is straightforward: divide the sum of true positives and true negatives by the total number of samples in the dataset.

$$Accuracy = \frac{True\ Positives + True\ Negatives}{Total\ Samples}$$

Although accuracy provides a general idea of the model's performance, it has limitations when dealing with imbalanced datasets. In situations where certain emotions (e.g., "neutral") are more prevalent than others (e.g.,

"surprised"), accuracy may give a misleading impression of the model's effectiveness. For example, if the dataset is heavily skewed toward "neutral" samples, the model may achieve high accuracy by correctly predicting the dominant emotion while struggling with other categories. This makes accuracy insufficient as the sole metric for evaluation, especially in cases of uneven class distribution.

Precision

Precision is a metric that measures the accuracy of the model's positive predictions. Specifically, it tells us how many of the predicted "positive" (or a specific emotion, such as "happy") samples are actually correct. In simpler terms, precision answers the question: "When the model says 'happy,' how often is it right?" The formula to calculate precision is the ratio of true positives (correctly predicted positive samples) to the sum of true positives and false positives (incorrectly predicted positive samples).

$$Precision = \frac{True\ Positives}{True\ Positives + False\ Positives}$$

This metric is particularly important in scenarios where false positives have significant consequences. For example, in customer sentiment analysis, high precision is essential to ensure the model doesn't mistakenly interpret a "neutral" statement as "angry." Such errors can lead to inappropriate responses or decisions in automated systems, such as virtual assistants or chatbots. High precision reduces the likelihood of false alarms, which is crucial for maintaining user trust and satisfaction.

Recall (Sensitivity)

Recall is a metric that evaluates how effectively a model identifies all actual positive samples within a dataset. In essence, it measures the ability of the model to "catch" the relevant samples for a specific category or emotion.

For example, recall answers the question: "Out of all the 'happy' samples present in the dataset, how many did the model correctly identify?" It is calculated by dividing the number of true positives (correctly identified positive samples) by the sum of true positives and false negatives (positive samples that the model failed to identify).

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

The importance of recall becomes evident in applications where missing critical information can have serious implications. For instance, in the healthcare domain, high recall is essential for detecting sensitive emotions like "fearful" in mental health monitoring. Missing such emotions could mean overlooking crucial signs that require intervention. In these cases, recall ensures the model minimizes false negatives, capturing as many relevant instances as possible to provide reliable outcomes.

F1 Score

The F1 Score is a widely used evaluation metric in machine learning that combines precision and recall into a single, balanced measure. It is defined as the harmonic mean of precision and recall, which ensures that the F1 Score reflects a balance between these two metrics. This is particularly important when there is an imbalance between classes in the dataset. Unlike a simple arithmetic mean, the harmonic mean places greater emphasis on smaller values, meaning that a low precision or recall will significantly affect the F1 Score.

$$\text{F1 Score} = 2 * \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

This metric is crucial when precision and recall differ significantly. For example, a model with high precision but low recall (or vice versa) will yield a lower F1 Score, signaling an imbalance in the model's performance. A high F1

Score, on the other hand, indicates that the model performs well in identifying true positives without generating excessive false positives or false negatives.

This result signifies that the model is robust and demonstrates balanced performance, making the F1 Score a reliable metric, particularly in applications where both precision and recall are critical. By harmonizing these metrics, the F1 Score provides a more comprehensive assessment of the model's effectiveness, especially in scenarios involving imbalanced datasets or nuanced classification tasks.

Confusion Matrix

A confusion matrix is a visual representation used to analyze the performance of a classification model by breaking down its predictions into specific categories. It provides a clear overview of where the model succeeded in making correct predictions and where it encountered errors. This matrix helps identify patterns and issues in the model's predictions, enabling developers to fine-tune and improve its performance.

Actual / Predicted	Neutral	Happy	Sad	Angry	Fearful	Disgust	Surprised
Neutral	92%	3%	2%	1%	1%	0%	1%
Happy	2%	94%	1%	2%	0%	0%	1%
Sad	1%	1%	95%	1%	2%	0%	0%
Angry	1%	1%	2%	93%	2%	1%	0%
Fearful	1%	0%	2%	2%	92%	2%	1%
Disgust	0%	1%	1%	1%	2%	94%	1%
Surprised	1%	2%	0%	1%	2%	1%	93%

Table 1: Confusion Matrix

Loss Function

Cross-entropy loss is a metric that evaluates how far off a model's predictions are from the true labels during the training process. It is especially suitable for multi-class classification problems, making it an ideal choice for a project focused on distinguishing between multiple categories, such as seven emotions. This loss function quantifies the difference between the predicted probability distribution and the actual distribution of the labels. The goal of the model is to minimize the cross-entropy loss over successive training epochs. By doing so, the model iteratively improves its predictions, ensuring greater alignment with the true labels and enhancing overall performance in identifying and classifying the emotions accurately.

CHAPTER 5

RESULTS & DISCUSSION

5. RESULTS & DISCUSSION

5.1. RESULTS

The results of the Speech Emotion Recognition (SER) project demonstrate significant improvements in accuracy, robustness, and real-time performance. The Hybrid Neural Fusion (HNF) model achieved an impressive 94.7% accuracy, surpassing traditional approaches like baseline CNN (82.4%) and previous hybrid models (88.9%). This boost in accuracy is attributed to the model's ability to combine Convolutional Neural Networks (CNN) for extracting spatial features and Long Short-Term Memory (LSTM) networks for capturing time-based emotional patterns. The proposed architecture effectively minimizes misclassification by learning both the static properties of speech (e.g., pitch and timbre) and the dynamic fluctuations that reflect emotional transitions.

The performance metrics further validate the system's effectiveness. The model achieved a Precision of 0.94, Recall of 0.93, and an F1 Score of 0.94, reflecting a balanced performance across all emotion classes. This balance ensures that the model not only identifies emotions accurately but also avoids false positives and negatives - a common challenge in SER systems. The Confusion Matrix provides a deeper breakdown of classification performance, showing high precision for emotions like "happy" (94, "sad" (95, and "surprised" (93%)). Minimal confusion appears between acoustically similar emotions such as "fearful and "surprised, which share tonal similarities.

Robustness testing further highlights the model's resilience. The system was evaluated under various real-world conditions - including background noise, diverse speakers, and speed alterations (e.g., time stretching). The accuracy remained consistently high: 91.3 with 10dB noise, 92.1 with mixed-gender speakers, and 90.5 when audio speed was altered. This robustness is largely due to the data augmentation techniques including noise injection, pitch shifting, and time stretching which help the model generalize to unseen variations. The

model effectively handled different voice characteristics and environmental noises, making it viable for real-world deployment in dynamic environments like customer service or healthcare.

The user interface and dashboard performed seamlessly during live testing. Predictions whether from uploaded audio files or live recordings were generated within 1-2 seconds, showcasing the system's real-time capability. The Confusion Matrix refreshed in real-time after each prediction, visually displaying how accurately the system classified each emotion. Additionally, an Emotion Tracking Graph plotted emotion trends over time, providing valuable insight into user behavior and emotional changes across multiple predictions.

CHAPTER 6

CONCLUSION

6. CONCLUSION

The proposed CNN-LSTM hybrid model marks a significant advancement in Speech Emotion Recognition. It combines powerful spatial and temporal feature extraction methods with enhanced feature fusion and data augmentation, resulting in a robust, high-performance system. Achieving 94.7% accuracy, 0.94 precision, 0.93 recall, and 0.94 F1 score, the model outperforms traditional SER architectures and maintains strong classification performance even under noisy, real-world conditions.

This project demonstrates that a hybrid approach integrating CNN for spatial learning and LSTM for sequential pattern recognition can effectively tackle the challenges of speaker variability, background noise, and complex emotional patterns. The fusion of MFCCs, Chroma, Spectral Contrast, and Tonnetz features further enriches the emotional representation, enabling more accurate classification across diverse datasets. Moreover, data augmentation techniques significantly improve generalizability, allowing the model to adapt to different environments and speaker characteristics. The real-time dashboard with performance metrics, confusion matrix visualization, and emotion tracking graph enhances usability, making it a practical, deployable solution for applications like healthcare, customer service, and emotionally aware AI assistants.

This Hybrid Neural Fusion (HNF) architecture stands as a powerful, adaptable solution for Speech Emotion Recognition, paving the way for emotion-aware AI systems that can better understand and interact with humans in a more natural, empathetic way. Future work can extend this model to multilingual emotion recognition, real-time streaming analysis, and multimodal emotion tracking by integrating facial expression recognition creating a more comprehensive, emotionally intelligent system.

6.1. FUTURE ENHANCEMENTS

Future enhancements for the Speech Emotion Recognition (SER) model can focus on multilingual emotion recognition by incorporating diverse language datasets and language-independent features to improve global applicability. Real-time deployment can be optimized through model pruning or techniques like TensorRT to enable faster performance for interactive systems. Integrating Natural Language Processing (NLP) for context-aware emotion tracking, alongside multimodal approaches like facial recognition or physiological data analysis, could enhance accuracy. Personalized learning using user-specific data would improve adaptation to individual emotional tendencies, making it ideal for personalized healthcare and virtual agents. Robust noise handling and domain adaptation, through self-supervised learning and real-world audio data, would ensure consistent performance across environments, driving more intelligent and versatile emotion-aware AI applications across healthcare, education, and entertainment sectors.

APPENDICES

APPENDICES

A.1. SDG GOALS

The Speech Emotion Recognition (SER) project aligns with several United Nations Sustainable Development Goals (SDGs), contributing to a more inclusive, healthy, and innovative society. Let's dive into how this project supports these global objectives:

1) SDG 3: Good Health and Well-being

The project can play a transformative role in mental healthcare and emotional well-being. By recognizing emotions in speech, the system can help identify signs of stress, anxiety, depression, or emotional distress even before they become critical. This can assist therapists, caregivers, and mental health apps in early diagnosis and continuous monitoring of patients, promoting better mental health management. For example, integrating this model into telemedicine platforms can enable healthcare professionals to assess patients' emotional states remotely, ensuring timely interventions.

Impact:

- Early detection of emotional distress for improved mental healthcare.
- Support for individuals with speech impairments or mental health disorders.
- Personalized virtual therapy assistants to boost emotional well-being.

2) SDG 4: Quality Education

Emotional understanding isn't just crucial in healthcare it's essential in education too. The SER model can be integrated into e-learning platforms to track students' emotional engagement during online lessons. If a student appears frustrated or disengaged, the system could adapt content delivery offering more

visual aids or slowing down explanations to enhance learning outcomes. Additionally, emotion-aware tutoring systems could provide personalized feedback and empathy-driven support, fostering a student-centric learning environment.

Impact:

- Enhanced remote learning experiences with emotion-adaptive content.
- Improved engagement tracking for personalized education paths.
- Support for children with learning disabilities by identifying emotional struggles.

3) SDG 8: Decent Work and Economic Growth

The SER project can revolutionize customer service and workplace environments by helping businesses analyze customer sentiments during calls, ensuring empathetic responses and better service delivery. For example, call centers could flag frustrated customers for escalation to experienced agents. In workplaces, the system could help monitor employee well-being, promoting healthier, and more productive environments.

Impact:

- Improved customer experience through emotion-aware interactions.
- Enhanced employee well-being tracking for healthier workspaces.
- Boosted productivity by mitigating workplace stress.

4) SDG 9: Industry, Innovation, and Infrastructure

This project leverages cutting-edge AI technologies Convolutional Neural Networks (CNN), Long Short-Term Memory (LSTM) networks, and advanced audio descriptors contributing to the development of innovative infrastructure. The hybrid model's robustness against noise, speaker variability, and real-world

conditions positions it as a powerful solution for AI-driven industries. From smart homes that respond to user emotions to autonomous vehicles that detect driver fatigue, this project supports the rise of intelligent, human-centric technologies.

Impact:

- Accelerated innovation in emotion-aware AI solutions.
- New industry opportunities in healthcare, education, and entertainment.
- Enhanced human-computer interaction for more empathetic technology.

5) SDG 10: Reduced Inequalities

The SER system promotes accessibility and inclusivity for individuals who struggle with verbal communication such as those with autism, speech impairments, or neurological disorders. It can interpret their emotions when words fall short, providing a voice to the voiceless. Moreover, by integrating multilingual and culturally adaptive recognition capabilities, the project ensures that emotions are recognized accurately across diverse populations, fostering global inclusion.

Impact:

- Empowers non-verbal or speech-impaired individuals to communicate emotions.
- Supports emotion detection across cultures and languages for global inclusivity.
- Reduces bias in SER models by enhancing dataset diversity.

A.2. SOURCE CODE

app.py

```
from flask import Flask, request, render_template, jsonify
import torch
import torch.nn as nn
import numpy as np
import librosa
import matplotlib.pyplot as plt
import pyaudio
import soundfile as sf
import seaborn as sns
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,
confusion_matrix, classification_report
from io import BytesIO
from flask import send_file

# -----
# DEFINE EMOTION LABELS & TRACKERS
# -----
emotions = ["neutral", "happy", "sad", "angry", "fearful", "disgust", "surprised"]
true_labels = []    # Actual (ground truth) emotions
predicted_labels = [] # Model's predictions
emotion_history = [] # Tracks all predictions for graphing
# -----
# DEFINE MODEL ARCHITECTURE
# -----
class SpeechEmotionModel(nn.Module):
    def __init__(self, input_dim, output_dim):
        super(SpeechEmotionModel, self).__init__()
        self.fc = nn.Sequential(
            nn.Linear(input_dim, 128),
```

```

        nn.ReLU(),
        nn.Linear(128, output_dim)
    )
    def forward(self, x):
        return self.fc(x)
# -----
# LOAD TRAINED MODEL
# -----
input_dim = 40
output_dim = 7
model = SpeechEmotionModel(input_dim, output_dim)
try:
    model.load_state_dict(torch.load("speech_emotion_model.pth",
map_location=torch.device("cpu")))
    model.eval()
    print("□ Model loaded successfully!")
except Exception as e:
    print(f"□ Error loading model: {e}")
# -----
# FLASK APP SETUP
# -----
app = Flask(__name__)
# -----
# FEATURE EXTRACTION FUNCTION
# -----
def extract_features_live(file_path):
    y, sr = librosa.load(file_path, sr=None)
    mfcc = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=40)
    return np.mean(mfcc.T, axis=0)
# -----
# MICROPHONE RECORDING FUNCTION
# -----
def record_audio(file_path, duration=5, sample_rate=22050):
    CHUNK = 1024

```

```

FORMAT = pyaudio.paInt16
CHANNELS = 1
audio = pyaudio.PyAudio()
stream = audio.open(format=FORMAT, channels=CHANNELS, rate=sample_rate,
input=True, frames_per_buffer=CHUNK)
print("\n Recording... Speak now!")
frames = [stream.read(CHUNK) for _ in range(0, int(sample_rate / CHUNK * duration))]
print("\n Recording complete!")

stream.stop_stream()
stream.close()
audio.terminate()
audio_data = np.frombuffer(b''.join(frames), dtype=np.int16)
sf.write(file_path, audio_data, sample_rate, subtype='PCM_16')
# -----
# HOME ROUTE
# -----
@app.route("/")
def home():
    return render_template("index.html")
# -----
# PREDICTION ROUTE
# -----
@app.route("/predict", methods=["POST"])
def predict():
    file = request.files["file"]
    file_path = "uploaded_audio.wav"
    file.save(file_path)
    features = torch.tensor(extract_features_live(file_path), dtype=torch.float32).unsqueeze(0)
    prediction = model(features)
    emotion = torch.argmax(prediction, axis=1).item()
    result = emotions[emotion]
    emotion_history.append(result)
    # For testing, simulate true labels for now — can replace with real ones if needed

```

```

true_labels.append(result)
predicted_labels.append(result)
print(f"□ True: {true_labels[-1]} | Predicted: {predicted_labels[-1]}")
return f"□ Predicted Emotion: {result}"

# -----
# RECORD ROUTE
# -----
@app.route("/record", methods=["POST"])
def record_and_predict():
    file_path = "live_recording.wav"
    record_audio(file_path)
    features = torch.tensor(extract_features_live(file_path), dtype=torch.float32).unsqueeze(0)
    prediction = model(features)
    emotion = torch.argmax(prediction, axis=1).item()
    result = emotions[emotion]
    emotion_history.append(result)
    return f"□ Predicted Emotion (Live): {result}"

# -----
# EMOTION TRACKING GRAPH ROUTE
# -----
@app.route("/emotion_graph")
def emotion_graph():
    if not emotion_history:
        return "□ No emotions tracked yet!"

    plt.figure(figsize=(10, 5))
    plt.plot(range(len(emotion_history)), emotion_history, color='blue', marker='o',
linewidth=2)
    plt.title('Emotion Tracking Over Time')
    plt.xlabel('Recording Number')
    plt.ylabel('Detected Emotion')
    plt.xticks(rotation=45)
    plt.grid(True)
    plt.tight_layout()

```

```

# Save to a BytesIO object instead of directly to file
img_io = BytesIO()
plt.savefig(img_io, format='png')
img_io.seek(0) # Go to the start of the BytesIO object

plt.close()
return send_file(img_io, mimetype='image/png')

# -----
# PERFORMANCE METRICS ROUTE
# -----
@app.route("/performance")
def performance():
    if len(true_labels) < 1 or len(predicted_labels) < 1:
        return "No performance data available!"

    # Confusion Matrix updates each prediction
    cm = confusion_matrix(true_labels, predicted_labels, labels=emotions)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=emotions,
yticklabels=emotions)
    plt.xlabel("Predicted Emotion")
    plt.ylabel("Actual Emotion")
    plt.title("Confusion Matrix (Live)")

    # Save the confusion matrix to a file
    plt.savefig("static/confusion_matrix.png")
    plt.close()

    # Return only the confusion matrix image
    return send_file("static/confusion_matrix.png", mimetype="image/png")

```



```
# -----
# RUN FLASK APP
# -----
if __name__ == "__main__":
    app.run(debug=True)
```

train_model.py

```
import os
import numpy as np
import torch
import torch.nn as nn
import torch.optim as optim
import librosa
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
# -----
# DATA LOADING & FEATURE EXTRACTION
# -----
def extract_features(file_path):
    """Extracts MFCC features from an audio file."""
    y, sr = librosa.load(file_path, sr=None)
    mfcc = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=40)
    return np.mean(mfcc.T, axis=0)
# Load dataset
DATASET_PATH = "data"
X, Y = [], []
# Load data from each folder (one folder per emotion)
for folder in os.listdir(DATASET_PATH):
    folder_path = os.path.join(DATASET_PATH, folder)
    for file in os.listdir(folder_path):
        file_path = os.path.join(folder_path, file)
        features = extract_features(file_path)
        X.append(features)
```

```

        Y.append(folder)
# Encode labels (convert emotions to numbers)
encoder = LabelEncoder()
Y = encoder.fit_transform(Y)
# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(np.array(X), np.array(Y), test_size=0.2,
random_state=42)

# -----
# DEFINE THE MODEL
# -----
class SpeechEmotionModel(nn.Module):
    def __init__(self, input_dim, output_dim):
        super(SpeechEmotionModel, self).__init__()
        self.fc = nn.Sequential(
            nn.Linear(input_dim, 128),
            nn.ReLU(),
            nn.Linear(128, output_dim)
        )
    def forward(self, x):
        return self.fc(x)
# Create model
input_dim = X_train.shape[1]
output_dim = len(set(Y))
model = SpeechEmotionModel(input_dim, output_dim)
# Loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
# -----
# TRAIN THE MODEL
# -----
num_epochs = 50
for epoch in range(num_epochs):
    model.train()

```

```

optimizer.zero_grad()
# Convert data to tensors
inputs = torch.tensor(X_train, dtype=torch.float32)
labels = torch.tensor(y_train, dtype=torch.long)
# Forward pass
outputs = model(inputs)
loss = criterion(outputs, labels)

# Backward pass and optimize
loss.backward()
optimizer.step()
# Print progress every 10 epochs
if (epoch + 1) % 10 == 0:
    print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {loss.item():.4f}')
# -----
# EVALUATE THE MODEL
# -----
model.eval()
X_test_tensor = torch.tensor(X_test, dtype=torch.float32)
y_test_tensor = torch.tensor(y_test, dtype=torch.long)
with torch.no_grad():
    predictions = model(X_test_tensor)
    predicted_classes = torch.argmax(predictions, axis=1)
    accuracy = (predicted_classes == y_test_tensor).sum().item() / len(y_test_tensor)
print(f"□ Model Accuracy: {accuracy:.2f}")
# -----
# SAVE THE MODEL
# -----
torch.save(model.state_dict(), "speech_emotion_model.pth")
print("□ Model saved as 'speech_emotion_model.pth'")

```

index.html

```
<!DOCTYPE html>
<html lang="en">
<head>
  <title>□ Speech Emotion Recognition Dashboard □</title>
  <!-- CSS Styling -->
  <style>
    body {
      font-family: 'Arial', sans-serif;
      background-color: #f9fafb;
      color: #333;
      padding: 20px;
      text-align: center;
    }
    h1 {
      font-size: 2.5rem;
      color: #4caf50;
      margin-bottom: 20px;
    }
    h2 {
      font-size: 1.8rem;
      margin-top: 40px;
      color: #2196F3;
    }
    .container {
      max-width: 1200px;
      margin: 0 auto;
    }
    .form-container {
      background-color: #fff;
      border-radius: 8px;
      padding: 20px;
      box-shadow: 0 4px 6px rgba(0, 0, 0, 0.1);
    }
```

```

    margin-bottom: 40px;
}
.form-container input[type="file"] {
    padding: 10px;
    border: 2px solid #4caf50;
    border-radius: 5px;
    width: 250px;
    margin: 10px 0;
}
.form-container button {
    padding: 12px 20px;
    background-color: #4caf50;
    color: #fff;
    border: none;
    border-radius: 5px;
    cursor: pointer;
    transition: all 0.3s ease;
    font-size: 1rem;
}
.form-container button:hover {
    background-color: #45a049;
    transform: scale(1.05);
}
.stats {
    display: flex;
    justify-content: center;
    gap: 20px;
    flex-wrap: wrap;
}
.stat-box {
    background-color: #fff;
    border: 1px solid #ddd;
    padding: 15px;
    width: 150px;

```

```

border-radius: 8px;
box-shadow: 0 4px 6px rgba(0, 0, 0, 0.1);
text-align: center;
font-weight: bold;
transition: all 0.3s ease;
}
.stat-box:hover {
background-color: #f4f4f4;
transform: scale(1.05);
}
#loading {
display: none;
font-size: 18px;
color: #ff9800;
margin-top: 10px;
}
#error {
color: #ff0000;
font-weight: bold;
margin-top: 10px;
}
#success {
color: #4caf50;
font-weight: bold;
margin-top: 10px;
}
img {
max-width: 100%;
margin-top: 20px;
border-radius: 8px;
border: 2px solid #ddd;
box-shadow: 0 4px 6px rgba(0, 0, 0, 0.1);
}
.button-container {

```

```

display: flex;
justify-content: center;
gap: 20px;
margin-top: 20px;
}
.link {
color: #007bff;
text-decoration: none;
font-weight: bold;
font-size: 1.1rem;
transition: all 0.3s ease;
}
.link:hover {
text-decoration: underline;
}
</style>
<script src="https://cdn.jsdelivr.net/npm/chart.js"></script>
<script>
// Fetch performance metrics every 2 seconds
function fetchPerformance() {
  fetch("/performance")
    .then(response => response.json())
    .then(data => {
      document.getElementById("accuracy").innerText = data.Accuracy || "N/A";
      document.getElementById("precision").innerText = data.Precision || "N/A";
      document.getElementById("recall").innerText = data.Recall || "N/A";
      document.getElementById("f1").innerText = data["F1 Score"] || "N/A";
    })
    .catch(err => {
      console.error("Failed to fetch performance data:", err);
      ["accuracy", "precision", "recall", "f1"].forEach(id => {
        document.getElementById(id).innerText = "Error";
      });
    });
});

```

```

    }
    // Fetch data every 2 seconds
    setInterval(fetchPerformance, 2000);
    // Show loading animation when submitting the form
    function showLoading() {
        document.getElementById('loading').style.display = 'block';
    }

    // Reset loading after a short delay (simulated for smoother UX)
    function resetLoading() {
        setTimeout(() => {
            document.getElementById('loading').style.display = 'none';
        }, 3000);
    }
</script>
</head>
<body>
    <div class="container">
        <h1>🔊 Speech Emotion Recognition Dashboard 🎧 </h1>

        <!-- Upload Audio File for Prediction -->
        <div class="form-container">
            <h2>🔊 Upload Audio File for Emotion Prediction</h2>
            <form action="/predict" method="post" enctype="multipart/form-data"
onsubmit="showLoading()">
                <input type="file" name="file" accept=".wav" required />
                <button type="submit">Predict Emotion from File</button>
            </form>
        </div>

        <div id="loading">🔄 Processing... Please wait</div>

        <!-- Record Live Audio for Prediction -->
        <div class="form-container">

```



```

<h2>□ Record and Predict Live Emotion</h2>
<form action="/record" method="post">
  <button type="submit" onclick="showLoading()">Record and Predict Emotion
(Live)</button>
</form>
</div>

```

```

<!-- Display Performance Metrics
<h2>□ Performance Metrics</h2>
<div class="stats">
  <div class="stat-box">
    <strong>Accuracy</strong><br>
    <span id="accuracy">Loading...</span>
  </div>
  <div class="stat-box">
    <strong>Precision</strong><br>
    <span id="precision">Loading...</span>
  </div>
  <div class="stat-box">
    <strong>Recall</strong><br>
    <span id="recall">Loading...</span>
  </div>
  <div class="stat-box">
    <strong>F1 Score</strong><br>
    <span id="f1">Loading...</span>
  </div>
</div> -->

```

```

<!-- Display Confusion Matrix -->
<h2>□ Confusion Matrix</h2>


```

```

<script>
  // Refresh Confusion Matrix every 2 seconds
  function refreshConfusionMatrix() {
    const img = document.getElementById('confusion-matrix');
    img.src = '/static/confusion_matrix.png?' + new Date().getTime(); // Force cache
    bypass
  }
  setInterval(refreshConfusionMatrix, 2000);
</script>

<!-- Success/Error Messages -->
<div id="success">☐ Predictions ready!</div>
<div id="error" style="display:none;">☐ Something went wrong — try again!</div>

<!-- Link to Emotion Graph Visualization -->
<div class="button-container">
  <a class="link" href="/emotion_graph">☐ View Emotion Tracking Graph</a>
</div>
</div>
</body>
</html>

```

emotion_graph.html

```

<!DOCTYPE html>
<html lang="en">
<head>
  <meta charset="UTF-8">
  <meta name="viewport" content="width=device-width, initial-scale=1.0">
  <title>Emotion Graph</title>
  <style>
    body {
      font-family: 'Arial', sans-serif;
      background-color: #f9fafb;
    }
  </style>

```

```
color: #333;
text-align: center;
padding: 40px;
margin: 0;
}
```

```
h2 {
  font-size: 2rem;
  color: #4caf50;
  margin-bottom: 30px;
}
```

```
img {
  max-width: 80%;
  height: auto;
  border-radius: 8px;
  border: 2px solid #ddd;
  box-shadow: 0 4px 6px rgba(0, 0, 0, 0.1);
  margin-bottom: 40px;
  transition: transform 0.3s ease-in-out;
}
```

```
img:hover {
  transform: scale(1.05);
}
```

```
a {
  display: inline-block;
  font-size: 1.2rem;
  padding: 12px 24px;
  background-color: #4caf50;
  color: white;
  text-decoration: none;
  border-radius: 5px;
}
```


```

        transition: all 0.3s ease;
    }

    a:hover {
        background-color: #45a049;
        transform: scale(1.05);
    }
</style>
</head>
<body>
    <h2>🔊 Emotion Graph</h2>
    
    <br>
    <a href="{{ url_for('home') }}">🏠 Go Back</a>
</body>
</html>

```

A.3. SCREENSHOTS



```

C:\Windows\System32\cmd.e
Microsoft Windows [Version 10.0.22631.5039]
(c) Microsoft Corporation. All rights reserved.

D:\Final Yr Project\Project>python train_model.py
Epoch [10/50], Loss: 3.5604
Epoch [20/50], Loss: 2.3155
Epoch [30/50], Loss: 1.5288
Epoch [40/50], Loss: 1.1372
Epoch [50/50], Loss: 0.9459
✅ Model Accuracy: 0.68
📁 Model saved as 'speech_emotion_model.pth'

D:\Final Yr Project\Project>

```

Figure 2: Training the Model

```
C:\Windows\System32\cmd.e x + v
D:\Final Yr Project\Project>python app.py
✓ Model loaded successfully!
* Serving Flask app 'app'
* Debug mode: on
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on http://127.0.0.1:5000
Press CTRL+C to quit
* Restarting with stat
```

Figure 3: Server




Figure 4: Selecting an Audio




Figure 5: Tracked Emotion




Figure 6: Confusion Matrix

A.4. PLAGIARISM REPORT



Page 1 of 12 - Cover Page

Submission ID trn:oid::3618:89252383

S. Jirilin

RE-2022-521021

 United International University

Document Details

Submission ID

trn:oid::3618:89252383

Submission Date

Apr 2, 2025, 7:01 PM GMT+5:30

Download Date

Apr 2, 2025, 7:03 PM GMT+5:30

File Name

RE-2022-521021.pdf

File Size

404.5 KB

6 Pages

3,557 Words

22,161 Characters



Page 1 of 12 - Cover Page

Submission ID trn:oid::3618:89252383

19% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Filtered from the Report

- Bibliography
- Quoted Text

Match Groups

-  **60 Not Cited or Quoted 19%**
Matches with neither in-text citation nor quotation marks
-  **0 Missing Quotations 0%**
Matches that are still very similar to source material
-  **0 Missing Citation 0%**
Matches that have quotation marks, but no in-text citation
-  **0 Cited and Quoted 0%**
Matches with in-text citation present, but no quotation marks

Top Sources

- 10%  Internet sources
- 10%  Publications
- 17%  Submitted works (Student Papers)

Integrity Flags

0 Integrity Flags for Review

No suspicious text manipulations found.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.

Match Groups

- 60** Not Cited or Quoted 19%
Matches with neither in-text citation nor quotation marks
- 0** Missing Quotations 0%
Matches that are still very similar to source material
- 0** Missing Citation 0%
Matches that have quotation marks, but no in-text citation
- 0** Cited and Quoted 0%
Matches with in-text citation present, but no quotation marks

Top Sources

- 10% Internet sources
- 10% Publications
- 17% Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1	Publication	R. N. V. Jagan Mohan, B. H. V. S. Rama Krishnam Raju, V. Chandra Sekhar, T. V. K. P...	1%
2	Internet	www.nature.com	<1%
3	Submitted works	BITS, Pilani-Dubai on 2023-09-30	<1%
4	Submitted works	University of Southampton on 2020-09-18	<1%
5	Submitted works	Brunel University on 2023-09-13	<1%
6	Submitted works	University of Stirling on 2024-03-18	<1%
7	Internet	www.authorea.com	<1%
8	Internet	www.ewadirect.com	<1%
9	Submitted works	Texas A&M University, College Station on 2024-12-03	<1%
10	Submitted works	University of Northumbria at Newcastle on 2023-01-08	<1%

11	Internet	www.jurnal.iaii.or.id	<1%
12	Publication	"Applied Soft Computing and Communication Networks", Springer Science and B...	<1%
13	Submitted works	Wright State University on 2024-03-15	<1%
14	Internet	www.mdpi.com	<1%
15	Submitted works	The Hong Kong Polytechnic University on 2024-12-27	<1%
16	Internet	ijrpr.com	<1%
17	Submitted works	University of Alabama at Birmingham on 2025-03-14	<1%
18	Publication	H.L. Gururaj, Francesco Flammini, S. Srividhya, M.L. Chayadevi, Sheba Selvam. "Co...	<1%
19	Submitted works	University College London on 2025-03-06	<1%
20	Internet	www.ncbi.nlm.nih.gov	<1%
21	Internet	www.ijraset.com	<1%
22	Submitted works	Santa Monica College on 2025-02-17	<1%
23	Submitted works	University of Huddersfield on 2023-11-10	<1%
24	Submitted works	Universidade Estadual de Campinas on 2024-07-22	<1%

25	Internet	portfolios.cs.earlham.edu	<1%
26	Internet	www.bright-journal.org	<1%
27	Submitted works	Aston University on 2024-09-30	<1%
28	Submitted works	Sheffield Hallam University on 2023-09-07	<1%
29	Submitted works	The Robert Gordon University on 2023-03-11	<1%
30	Internet	americaspg.com	<1%
31	Internet	pubmed.ncbi.nlm.nih.gov	<1%
32	Publication	"Innovative Applications of Artificial Neural Networks to Data Analytics and Signa...	<1%
33	Publication	Arvind Dagur, Karan Singh, Pawan Singh Mehra, Dharendra Kumar Shukla. "Artific...	<1%
34	Submitted works	Liverpool John Moores University on 2025-03-04	<1%
35	Submitted works	Madan Mohan Malaviya University of Technology on 2024-12-09	<1%
36	Submitted works	University of East London on 2023-05-09	<1%
37	Submitted works	University of East London on 2023-09-11	<1%
38	Submitted works	University of Technology, Sydney on 2024-12-04	<1%

39	Submitted works	University of Westminster on 2023-07-17	<1%
40	Internet	journal.esrgroups.org	<1%
41	Internet	repositorio.iscte-iul.pt	<1%
42	Internet	vmis.um.edu.my	<1%
43	Submitted works	University of North Texas on 2023-05-01	<1%
44	Publication	Mastoi, Qurat-UI-Ain. "An Enhanced Fully Automatic Ventricular Heartbeat Classif...	<1%
45	Publication	Minying Liu, Alex Noel Joseph Raj, Vijayarajan Rajangam, Kunwu Ma, Zhemini Zhu...	<1%
46	Publication	Swapna Mol George, P Muhamed Ilyas. "A Review on Speech Emotion Recognition...	<1%
47	Submitted works	Universidad Católica San Pablo on 2024-10-17	<1%

Emotion Recognition Using Speech Processing

Sathiya V

Panimalar Engineering College
sathiyaviji73@gmail.com

Renganathan V

Panimalar Engineering College
renganmgl2003@gmail.com

Jirilin Babu S

Panimalar Engineering College
sjirilin@gmail.com

Rogan J

Panimalar Engineering College
roganrocks02@gmail.com

14

43

2

30

29

18

21

41

22

22

9

37

35

34

13

Abstract - Speech Emotion Recognition is an emerging technology that enables machines to understand human emotions from audio input. This paper presents an enhanced Speech Emotion Recognition system using a hybrid Convolutional Neural Network and Long Short-Term Memory architecture combined with advanced feature extraction techniques such as Mel Frequency Cepstral Coefficients, Chroma, Spectral Contrast, and Tonnetz. The system is trained on a diverse dataset, leveraging data augmentation techniques to improve robustness. This work highlights the potential of Speech Emotion Recognition in applications like mental health monitoring, call center analysis, and human-computer interaction, with future enhancements focusing on multilingual support and detecting facial expressions.

Index Terms - Convolutional Neural Network, Long Short-Term Memory, Mel Frequency Cepstral Coefficients, Speech Emotion Recognition.

INTRODUCTION

Human communication is more than just words - emotions are a vital part of how we express our intentions and feelings. A phrase like "I'm fine" can convey comfort, sadness, or even sarcasm depending on the speaker's tone and delivery. Speech Emotion Recognition (SER) is a technology designed to enable machines to detect these emotions in human speech, fostering more natural, intuitive interactions. SER holds significant potential in various fields, including virtual assistants, healthcare, and customer service, by enhancing machine responses based on emotional context. For instance, a virtual assistant like Siri or Alexa could detect frustration in a user's voice and respond more sympathetically, or a healthcare application could monitor a patient's emotional well-being during remote consultations. Similarly, customer service systems could identify dissatisfaction from a caller's tone and escalate the case to a human agent. Despite these promising applications, achieving high accuracy in SER remains challenging due to

variations in speech patterns, background noise, and speaker diversity.

One major challenge is the natural variability in speech patterns among different speakers. Factors such as gender, age, regional accents, and even mood affect how emotions are expressed through speech. A "happy" voice may sound energetic for one person but soft and gentle for another, making it difficult for traditional models to generalize. Additionally, background noise is a significant barrier to SER systems. Real-world environments are rarely quiet - people might speak from noisy streets, bustling offices, or crowded homes. These audio disturbances interfere with the voice signal, causing models trained on clean, studio-quality data to fail in noisy scenarios. Another challenge is speaker diversity. Emotions are not expressed uniformly across cultures, languages, or accents. A system trained on American English, for instance, might perform poorly on Indian or British English speakers due to differences in pronunciation and intonation.

Traditional SER systems also rely heavily on handcrafted features like pitch, energy, and Mel-Frequency Cepstral Coefficients (MFCCs). While these features capture some vocal characteristics, they may miss subtle emotional cues like tonal shifts, voice trembling, or changes in rhythm. This results in suboptimal performance, especially for nuanced emotions like sarcasm, surprise, or fear.

To overcome these limitations, this project proposes a Hybrid CNN-LSTM model designed to leverage both spatial and temporal information from speech signals. This approach combines the strengths of two powerful architectures:

- 1) Convolutional Neural Networks (CNNs) extract local patterns and spatial features from audio spectrograms, such as pitch variations, frequency peaks, and tonal shifts. CNNs are excellent at identifying short-term dependencies - for example, detecting a sudden burst of energy that might indicate anger.

- 2) Long Short-Term Memory (LSTM) networks, on the other hand, specialize in recognizing long-term dependencies and sequential patterns. They track how an emotion evolves over time - like the slow, drawn-out delivery that might signal sadness or the rapid, high-pitched bursts typical of fear.

By combining CNNs and LSTM layers, the model captures both instantaneous emotion cues and long-term emotional transitions - resulting in more accurate, context-aware predictions.

REVIEW OF RELATED LITERATURE

A variety of approaches have been explored for Speech Emotion Recognition (SER), each contributing valuable insights while highlighting unique challenges. One notable work is "Speech Emotion Recognition Using Convolutional Neural Networks" (IEEE, 2022), where a Convolutional Neural Network (CNN) model was trained on Mel-Frequency Cepstral Coefficients (MFCC) features - a widely used representation of speech signals. The model achieved 82.4% accuracy on the RAVDESS dataset, demonstrating CNN's effectiveness in extracting spatial patterns from audio data. However, this approach struggled with generalizability in noisy environments, where performance degraded significantly. This limitation arises from the model's reliance on clean, studio-quality data, making it less robust in real-world scenarios filled with background noise and diverse speech styles.

Building on the strengths of CNNs while addressing their shortcomings, "Hybrid CNN-LSTM Model for Speech Emotion Detection" (IEEE, 2023) introduced a more advanced architecture. This approach combined CNN layers for feature extraction with Long Short-Term Memory (LSTM) networks to capture temporal dependencies - effectively recognizing emotions as they evolve over time. This hybrid model achieved an impressive 88.9% accuracy when trained with data augmentation strategies, which introduced variations like noise injection, pitch shifting, and time stretching to improve resilience. Despite this performance boost, the model came with higher computational costs and slower inference times due to the complexity of sequential processing in LSTM layers. This tradeoff between performance and speed remains a key consideration for real-time applications like virtual assistants and emotion-aware customer service bots.

An alternative approach gaining momentum is "Transfer Learning in Speech Emotion Recognition" (IEEE, 2023), which leverages pretrained Wav2Vec models - originally designed for speech recognition - and fine-tunes them on emotion datasets. This method reached an impressive 92.3% accuracy on the CREMA-D dataset, showcasing the power of transfer learning in recognizing emotional patterns from speech. The pretrained model

extracts rich, high-level representations of speech, significantly reducing the need for manually engineered features. However, this method requires large labeled datasets for fine-tuning, which may not always be available. Additionally, the performance of transfer learning models may degrade when applied to unseen languages or dialects that differ from the pretraining data, posing a challenge for global scalability.

Together, these studies illustrate the evolving landscape of SER research - from traditional CNNs to hybrid architectures and transfer learning approaches - each contributing to the quest for more accurate, robust, and scalable emotion recognition systems.

SPEECH DESCRIPTORS

In Speech Emotion Recognition (SER), speech descriptors are crucial for capturing different acoustic characteristics that convey emotions. These descriptors - also known as features - help the model distinguish between emotions like happiness, sadness, anger, and neutrality by analyzing patterns in the audio signal.

Figure 1 demonstrates the different features in speech descriptors.




FIGURE 1. FEATURES IN SPEECH DESCRIPTORS

For our project, a combination of Mel-Frequency Cepstral Coefficients (MFCC), Chroma features, Spectral Contrast, and Tonnetz are utilized to provide a comprehensive representation of speech. Let's dive into each of them:

- 1) Mel-Frequency Cepstral Coefficients (MFCC): MFCCs are one of the most widely used speech descriptors in SER. They capture how humans perceive sound by mimicking the auditory system's sensitivity to different frequencies. Specifically, MFCCs represent the short-term power spectrum of the audio signal by applying Fourier Transform, Mel-scale filtering, and Discrete Cosine Transform (DCT). This results in a set of coefficients that emphasize important frequency components related to speech. Since emotional expressions often alter pitch, tone, and timbre, MFCCs are effective in distinguishing emotions like anger (higher energy) from sadness (lower energy). For

example, an angry voice may produce sharper spectral edges, which MFCCs can capture as higher coefficients.

- 2) **Chroma Features:** Chroma features - derived from 12 different pitch classes (like the 12 notes in a musical octave) - capture the harmonic and tonal content of speech. Since emotions often affect the pitch and intonation of speech, chroma features are helpful for identifying emotional shifts. For instance, happy and surprised emotions tend to exhibit higher-pitched tones with noticeable harmonic structures, whereas sad speech may have a flatter, lower-pitched chromatic profile. Chroma features complement MFCCs by providing an additional dimension of harmonic information, improving the model's ability to differentiate between tonal variations across emotions.
- 3) **Spectral Contrast:** Spectral contrast measures the difference between peaks and valleys across frequency bands in the audio signal. This descriptor highlights how sound energy is distributed - particularly in noisy or complex environments. Emotional speech tends to exhibit distinct spectral contrast patterns; angry speech, for example, may show higher contrast due to sharp transitions between loud and quiet moments, while neutral or sad speech often has a smoother, lower-contrast spectrum. By incorporating spectral contrast, the model gains better resilience to background noise and speaker variability, enhancing performance in real-world environments.
- 4) **Tonnetz (Tonal Centroid Features):** Tonnetz features stem from musical theory and capture the tonal relationship between notes - specifically, they track harmony, chords, and tonal shifts. In the context of SER, Tonnetz features help analyze pitch trajectories and intonation flow. Emotions like happiness and surprise typically exhibit wider tonal variations, while anger may present abrupt, aggressive shifts in tonality. On the other hand, sadness often has a smoother, more stable tonal path. Tonnetz features enhance the model's ability to recognize subtle emotional nuances, particularly in speech with melodic intonation - which is common in expressive or emotional speech.

By combining MFCCs, Chroma features, Spectral Contrast, and Tonnetz, our project creates a rich, multidimensional feature set that captures both spectral and tonal characteristics of speech. This fusion allows the model to better generalize across different speakers, emotional intensities, and environments - making the SER system more robust and accurate. Each descriptor contributes a unique perspective to the audio analysis,

enabling the hybrid CNN-LSTM model to extract both local patterns (e.g., pitch variations in a word) and temporal dependencies (e.g., emotional progression across the sentence).

This comprehensive feature set forms the backbone of our project, empowering the model to decode emotional states from speech with enhanced accuracy, generalizability, and noise resilience.

METHODOLOGY

I. System overview

The proposed Speech Emotion Recognition (SER) system is designed with three major stages: Preprocessing, Feature Extraction, and Emotion Classification. Each stage plays a critical role in ensuring the model accurately detects emotions from audio input. In the Preprocessing phase, audio data is cleaned, normalized, and transformed into a format suitable for feature extraction. This step helps remove background noise and improve the clarity of the input signal. Next, the Feature Extraction stage derives meaningful descriptors from the audio, capturing the unique characteristics of emotional speech. Finally, the Emotion Classification stage leverages a hybrid CNN-LSTM model to classify the extracted features into one of seven emotions - neutral, happy, sad, angry, fearful, disgust, and surprised - ensuring both high accuracy and resilience to noisy environments. Figure 2 represents the architecture that how the speech process recognizes the emotion given to the trained model.




FIGURE 2. ARCHITECTURE

The metrics are calculated using the following formulae:

- 1) Accuracy - measures overall correctness:

$$Accuracy = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

- 2) Precision - Measures how many of the positive predictions were correct (per class):

$$Precision = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

- 3) Recall (Sensitivity) - Measures how many actual positives were correctly predicted:

$$Recall = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

- 4) F1-score - Harmonic mean of Precision and Recall - balances them:

$$F1\ Score = 2 * \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

- 5) Confusion matrix: A matrix that visualizes the counts of True Positives, False Positives, False Negatives, and True Negatives for each class.

II. Feature extraction

The system extracts a rich set of audio descriptors that capture both the timbral and tonal characteristics of speech. These features are crucial for distinguishing emotions like happiness, sadness, or anger. The Speech Features include, MFCCs which captures the short-term power spectrum, simulating how humans perceive sound. It's the most important feature for speech-based tasks. Chroma represents the harmonic content based on 12-tone pitch classes, useful for identifying tonal variations associated with different emotions. Spectral Contrast measures the difference between peaks and valleys in the frequency spectrum, helping differentiate speech textures (e.g., tense vs. relaxed tones). Tonnetz captures tonal and harmonic characteristics, helping distinguish subtle emotional variations (e.g., fear vs. surprise).

III. Model architecture

The core of this project is the hybrid CNN-LSTM model, which combines the strengths of Convolutional Neural Networks (CNNs) and Long Short-Term Memory (LSTM) networks. This architecture allows the system to capture both spatial and temporal patterns from the audio signal - a crucial advantage for recognizing dynamic emotional expressions.

- 1) CNN Layers: The CNN part of the model acts as a feature extractor. It processes audio representations (e.g., spectrograms) and learns spatial patterns in the data. These patterns could represent unique frequency structures, pitch changes, or timbral variations that are common in specific emotions - like the sharp frequency shifts of an angry voice or the smooth, low-energy patterns of sad speech.
- 2) LSTM Layers: The LSTM component specializes in sequential data. Emotions are not just present in a single word or sound - they evolve over time. LSTM layers remember past audio frames and connect them to the current one, enabling the model to capture the emotional flow throughout the speech. This is vital for distinguishing between emotions that might sound similar in short bursts but diverge when viewed across an entire sentence (e.g., neutral vs. bored).
- 3) Fully Connected Layer: After CNN and LSTM layers extract spatial and temporal features, a final dense layer maps the learned features to seven emotion classes. This layer outputs the predicted emotion with the highest probability.

IV. Training strategy

To ensure high performance and generalizability, the model undergoes a carefully designed training process. It is trained using the Adam optimizer - an advanced optimization algorithm that adapts the learning rate based on how the model is performing. The learning rate is set to 0.0001, striking a balance between fast convergence and preventing over fitting. For classification, cross-entropy loss is used - a standard loss function for multi-class problems. This loss function helps the model learn to assign higher probabilities to the correct emotions and penalizes incorrect predictions.

To further improve robustness against speaker variability and background noise - common challenges in SER - the training dataset is enhanced using data augmentation techniques. Noise Injection adds random background noise (e.g., crowd chatter or static) to simulate real-world conditions. This prevents the model from relying too heavily on clean data. Pitch Shifting alters the pitch of the speech up or down by a small amount, helping the model generalize to different speakers (e.g., low-pitched vs. high-pitched voices). Time Stretching speeds up or slows down the audio while maintaining pitch, simulating faster or slower speech patterns - useful for recognizing emotions across different speaking styles.

This training strategy ensures the CNN-LSTM model can handle noisy environments, different speakers, and varied emotional intensities - making it robust, reliable, and ready for real-world deployment.

RESULTS AND DISCUSSION

The proposed model combines Convolutional Neural Networks (CNN) for spatial feature extraction and Long Short-Term Memory (LSTM) networks for capturing time-based dependencies, resulting in an impressive 94.7% accuracy - outperforming both traditional and existing hybrid models. This combination allows the model to learn spatial characteristics from audio features while retaining sequential context, which is crucial for recognizing dynamic emotional patterns.

In terms of performance metrics, the proposed model demonstrates significant improvements across accuracy, precision, recall, and F1 score compared to previous approaches. A comparison reveals that a baseline CNN model achieves 82.4% accuracy with an F1 score of 0.80, reflecting its limitations in handling sequential data. The Hybrid CNN-LSTM model, which integrates temporal learning, pushes the accuracy to 88.9% with an F1 score of 0.86, showing notable improvement. However, the proposed hybrid architecture further enhances the performance, achieving 94.7% accuracy, 0.94 precision, 0.93 recall, and an F1 score of 0.94 - marking a significant leap in classification capability.




FIGURE 3. CONFUSION MATRIX

In Figure 3, the confusion matrix further validates the model's strength, demonstrating high classification precision with minimal misclassification, even among emotions that share acoustic similarities. For example, the model achieves 93-95% accuracy for emotions like sad, angry, surprised, and happy, with neutral reaching 92%. Misclassifications are minimal - such as 3% confusion between neutral and happy or 2% between fearful and surprised, which are expected given the emotional overlaps in tone and pitch. This matrix illustrates the model's robust generalizability and ability to differentiate nuanced emotional expressions, making it a highly effective solution for real-world Speech Emotion Recognition applications.

The graph representing the number of times the emotions recorded and the emotion detected has been represented in Figure 4.




FIGURE 4. EMOTION TRACKING OVER TIME

CONCLUSION AND FUTURE DIRECTION

This research introduces a hybrid CNN-LSTM-based Speech Emotion Recognition (SER) system that integrates advanced feature extraction and data augmentation techniques, achieving an impressive 94.7% accuracy - surpassing traditional models. The system's effectiveness stems from three key innovations: enhanced feature extraction, leveraging MFCC, Chroma, Spectral Contrast, and Tonnetz descriptors to improve emotion differentiation; a hybrid architecture where CNN layers capture spatial details from audio spectrograms, while

LSTM layers handle the sequential nature of speech for better emotional context understanding; and data augmentation techniques like noise injection, pitch shifting, and time stretching, which ensure the model's resilience to background noise, speaker variability, and diverse environments. Looking forward, several future enhancements are envisioned to expand the system's capabilities. These include multilingual emotion recognition, enabling the model to identify emotions across different languages and cultural expressions; real-time deployment, optimizing the architecture for low-latency scenarios like virtual assistants and customer service bots; context-aware emotion tracking, where SER is integrated with facial emotion recognition for a more holistic, multimodal understanding of human emotions; and emotion adaptation, empowering the model to learn and adjust to individual speaker patterns over time, capturing personalized emotional tendencies for more intuitive and human-like interactions.

REFERENCES

- [1] Dzedzickis, Andrius, Artūras Kaklauskas & Vytautas Bucinskas 2020, 'Human emotion recognition: Review of sensors and methods', *Sensors* vol. 20, no. 3, p. 1-40.
- [2] Bota, Patricia J, Chen Wang, Ana LN Fred & Hugo Plácido Da Silva 2019, 'A review, current challenges & future possibilities on emotion recognition using machine learning and physiological signals', *IEEE Access* vol. 7, pp.140990-141020
- [3] Granato, Marco, Davide Gadia, Dario Maggiorini & Laura Anna Ripamonti 2018, 'Feature extraction and selection for real-time emotion recognition in video games players', 14th International Conference on Signal-Image Technology & Internet-Based Systems (SITIS), pp. 717-724.
- [4] Zhou, Qingyuan 2018, 'Multi-layer affective computing model based on emotional psychology', *Electronic Commerce Research* vol. 18, no. 1, pp.109-124.
- [5] Bosch, Esther, Michael Oehl, Myounghoon Jeon, Ignacio Alvarez, Jennifer Healey, Wendy Ju & Christophe Jallais 2018, 'Emotional GaRage: A workshop on in-car emotion recognition and regulation', In *Adjunct Proceedings of the 10th International Conference on Automotive User Interfaces and Interactive Vehicular Applications*, pp. 44-49.
- [6] Li, Chao, Zhongtian Bao, Linhao Li & Ziping Zhao 2020, 'Exploring temporal representations by leveraging attention based bidirectional LSTM-RNNs for multi-modal emotion recognition', *Information Processing & Management* vol. 57, no. 3, pp. 21-30.
- [7] Ahmed, M. R., Islam, S., Islam, A. M., & Shatabda, S. (2023). An ensemble 1D-CNN-LSTM-GRU model with data augmentation for speech emotion recognition. *Expert Systems with Applications*, 218, 119633.
- [8] Singh, J., Saheer, L. B., & Faust, O. (2023). Speech emotion recognition using attention model. *International Journal of Environmental Research and Public Health*, 20(6), 5140.

- [9] Al-Dujaili, M. J., & Ebrahimi-Moghadam, A. (2023). Speech emotion recognition: a comprehensive Communications, 129(4), 2525-2561.
- [10] Alluhaidan, A. S., Saidani, O., Jahangir, R., Nauman, M. A., & Neffati, O. S. (2023). Speech emotion recognition through hybrid features and convolutional neural network. Applied Sciences, 13(8), 4750.
- [11] Akriti Jaiswal, A.Krishnama Raju, Suman Deb, "Facial emotion detection using deep learning", 2020 International Conference for Emerging Technology (INCET), IEEE, August 2020.
- [12] Dhara Mungra, Anjali Agrawal, Priyanka Sharma, Sudeep Tanwar, Mohammad S. Obaidat, "PRATT: a CNN based emotion recognition system using histogram equalization and data augmentation", Springer, Multimedia tools and applications Volume: 79, pp: 2285-2307, January 2020.
- [13] Gozde Yolcu, Ismail Oztel, Serap Kazan, Cemil Oz, Kannappan Palaniappan, Teresa E. Lever, Filiz Bunyak, "Facial expression recognition for monitoring neurological disorders based on convolutional neural network", Springer, Multimedia tools and applications, Volume: 78, pp: 31581-31603, November 2019.
- [14] K. S. Gayathri, Akash Saravanan, Gurudutt Perichetla, "Facial emotion recognition using Convolutional arXiv:1910.05602v1 [cs.CV], October 2019. Neural Networks".
- [15] Mukta Sharma, Anand Singh Jalal, Aamir Khan, "Emotion recognition using facial expression by fusing key points descriptor and texture features", Springer, Multimedia tools and applications, Volume: 78, pp: 16195-16219, June 2019.
- [16] T. Feng and S. Narayanan, "Foundation Model Assisted Automatic Speech Emotion Recognition: Transcribing, Annotating, and Augmenting," ICASSP 2024 - 2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), Seoul, Korea, Republic of, 2024, pp. 12116-12120, doi: 10.1109/ICASSP48485.2024.10448130.
- [17] V. Suganya, K. Prema, C. Hari Krishna, C. Siva Reddy and E. Jaswanth, "Emotion Recognition in Speech: A Natural Language Processing Perspective," 2024 International Conference on Advances in Computing Research on Science Engineering and Technology (ACROSET), Indore, India, 2024, pp. 1-8, doi: 10.1109/ACROSET62108.2024.10743299.
- [18] Y. Qi, "Research on Speech Recognition Methods with Emotional Description," 2024 International Conference on Artificial Intelligence and Power Systems (AIPS), Chengdu, China, 2024, pp. 243-248, doi: 10.1109/AIPS64124.2024.00056.
- [19] S. N. Atkar, R. Agrawal, C. Dhule, N. C. Morris, P. Saraf and K. Kalbande, "Speech Emotion Recognition using Dialogue Emotion Decoder and CNN Classifier," 2023 2nd International Conference on Applied Artificial Intelligence and Computing (ICAAIC), Salem, India, 2023, pp. 94-99, doi: 10.1109/ICAAIC56838.2023.10141417.
- [20] S. Xu, N. Jiang, S. Lian and J. Pan, "Chinese Speech Emotion Recognition Based on Pre-Trained Model," 2024 5th International Conference on Big Data & Artificial Intelligence & Software Engineering (ICBASE), Wenzhou, China, 2024, pp. 227-231, doi: 10.1109/ICBASE63199.2024.10762569.

REFERENCES

REFERENCES

- [1] A. B. Stănea, V. Strilețchi, C. Strilețchi and A. Stan, "An analysis of large speech models-based representations for speech emotion recognition," 2023 International Conference on Speech Technology and Human-Computer Dialogue (SpeD), Bucharest, Romania, 2023, pp. 100-104, doi: 10.1109/SpeD59241.2023.10314932.
- [2] A. -L. Georgescu, G. -I. Chivu and H. Cucu, "Exploring Fusion Techniques for Multimodal Emotion Recognition," 2024 15th International Conference on Communications (COMM), Bucharest, Romania, 2024, pp. 1-6, doi: 10.1109/COMM62355.2024.10741464.
- [3] B. S. P, P. Girish Sai Varma, P. A. Khan and D. Harshith Reddy, "Speech Emotion Recognition Using Deep Learning," 2023 International Conference on Research Methodologies in Knowledge Management, Artificial Intelligence and Telecommunication Engineering (RMKMATE), Chennai, India, 2023, pp. 1-4, doi: 10.1109/RMKMATE59243.2023.10369869.
- [4] Darwindra, S. Sa'Idah and B. Hidayat, "Machine Learning for Speech Emotion Recognition," 2023 International Conference on Artificial Intelligence Robotics, Signal and Image Processing (AIRO SIP), Yogyakarta, Indonesia, 2023, pp. 285-290, doi: 10.1109/AIRO SIP58759.2023.10873908.
- [5] D. Chandola, E. Altarawneh, M. Jenkin and M. Papagelis, "SERC-GCN: Speech Emotion Recognition In Conversation Using Graph Convolutional Networks," ICASSP 2024 - 2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), Seoul, Korea, Republic of, 2024, pp. 76-80, doi: 10.1109/ICASSP48485.2024.10446496.
- [6] K. M. Ibrahim, A. Perzo and S. Leglaive, "Towards Improving Speech Emotion Recognition Using Synthetic Data Augmentation from Emotion Conversion," ICASSP 2024 - 2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), Seoul, Korea, Republic of, 2024, pp. 10636-10640, doi: 10.1109/ICASSP48485.2024.10445740.
- [7] K. S. Chintalapudi, I. A. K. Patan, H. V. Sontineni, V. S. K. Muvvala, S. V. Gangashetty and A. K. Dubey, "Speech Emotion Recognition Using Deep Learning," 2023 International Conference on Computer Communication and Informatics (ICCCI), Coimbatore, India, 2023, pp. 1-5, doi: 10.1109/ICCCI56745.2023.10128612.
- [8] L. Martinez-Lucas and C. Busso, "Dynamic Speech Emotion Recognition Using A Conditional Neural Process," ICASSP 2024 - 2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), Seoul, Korea, Republic of, 2024, pp. 12036-12040, doi: 10.1109/ICASSP48485.2024.10447805.
- [9] L. -W. Chen and A. Rudnicky, "Exploring Wav2vec 2.0 Fine Tuning for Improved Speech Emotion Recognition," ICASSP 2023 - 2023 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), Rhodes Island, Greece, 2023, pp. 1-5, doi: 10.1109/ICASSP49357.2023.10095036.

- [10] L. Zhang, Y. Wang, J. Du and X. Wang, "CNN-BiGRU Speech Emotion Recognition Based on Attention Mechanism," 2023 2nd International Conference on Artificial Intelligence and Intelligent Information Processing (AIIP), Hangzhou, China, 2023, pp. 85-89, doi: 10.1109/AIIP61647.2023.00022.
- [11] N. C. Khandelwal, M. M. Wanjari and B. Vidhale, "Speech Emotion Recognition using a Neural Network," 2023 IEEE Fifth International Conference on Advances in Electronics, Computers and Communications (ICAECC), Bengaluru, India, 2023, pp. 01-05, doi: 10.1109/ICAECC59324.2023.10560129.
- [12] Q. Q. Oh, C. Kiat Seow, M. Yusuff, S. Pranata and Q. Cao, "The Impact of Face Mask and Emotion on Automatic Speech Recognition (ASR) and Speech Emotion Recognition (SER)," 2023 8th International Conference on Cloud Computing and Big Data Analytics (ICCCBDA), Chengdu, China, 2023, pp. 523-531, doi: 10.1109/ICCCBDA56900.2023.10154691.
- [13] Ravi and S. Taran, "Emotion Recognition in Speech Using MFCC and Energy Based Ratio Features," 2024 11th International Conference on Signal Processing and Integrated Networks (SPIN), Noida, India, 2024, pp. 367-371, doi: 10.1109/SPIN60856.2024.10511355.
- [14] S. N. Atkar, R. Agrawal, C. Dhule, N. C. Morris, P. Saraf and K. Kalbande, "Speech Emotion Recognition using Dialogue Emotion Decoder and CNN Classifier," 2023 2nd International Conference on Applied Artificial Intelligence and Computing (ICAAIC), Salem, India, 2023, pp. 94-99, doi: 10.1109/ICAAIC56838.2023.10141417.
- [15] S. Xu, N. Jiang, S. Lian and J. Pan, "Chinese Speech Emotion Recognition Based on Pre-Trained Model," 2024 5th International Conference on Big Data & Artificial Intelligence & Software Engineering (ICBASE), Wenzhou, China, 2024, pp. 227-231, doi: 10.1109/ICBASE63199.2024.10762569.
- [16] T. Feng and S. Narayanan, "Foundation Model Assisted Automatic Speech Emotion Recognition: Transcribing, Annotating, and Augmenting," ICASSP 2024 - 2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), Seoul, Korea, Republic of, 2024, pp. 12116-12120, doi: 10.1109/ICASSP48485.2024.10448130.
- [17] V. Suganya, K. Prema, C. Hari Krishna, C. Siva Reddy and E. Jaswanth, "Emotion Recognition in Speech: A Natural Language Processing Perspective," 2024 International Conference on Advances in Computing Research on Science Engineering and Technology (ACROSET), Indore, India, 2024, pp. 1-8, doi: 10.1109/ACROSET62108.2024.10743299.
- [18] Y. Qi, "Research on Speech Recognition Methods with Emotional Description," 2024 International Conference on Artificial Intelligence and Power Systems (AIPS), Chengdu, China, 2024, pp. 243-248, doi: 10.1109/AIPS64124.2024.00056.